

Income_Predictor

September 29, 2020

1 Income Classifier

This Projects involves using census data from 1994 to explore if there is a way to accuratley predict if someone makes below or above 50k a year. In total there are 14 independent variables and roughly 45,000 data points between train and test sets. This project showcases preprocessing, data exploring, modeling, and model interpretation utilizing PyCaret, seaborn, and sklearn

```
[96]: from pycaret.classification import *
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import norm
from sklearn.preprocessing import StandardScaler
from scipy import stats
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
import sys
```

```
[14]: #Read in Data
data = pd.read_csv('au_train.csv')
data.head()
```

```
[14]:
```

	age	workclass	fnlwgt	education	education-num	\
0	39	State-gov	77516	Bachelors	13	
1	50	Self-emp-not-inc	83311	Bachelors	13	
2	38	Private	215646	HS-grad	9	
3	53	Private	234721	11th	7	
4	28	Private	338409	Bachelors	13	

	marital-status	occupation	relationship	race	sex	\
0	Never-married	Adm-clerical	Not-in-family	White	Male	
1	Married-civ-spouse	Exec-managerial	Husband	White	Male	
2	Divorced	Handlers-cleaners	Not-in-family	White	Male	
3	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	
4	Married-civ-spouse	Prof-specialty	Wife	Black	Female	

	capital-gain	capital-loss	hours-per-week	native-country	class
--	--------------	--------------	----------------	----------------	-------

0	2174	0	40	United-States	0
1	0	0	13	United-States	0
2	0	0	40	United-States	0
3	0	0	40	United-States	0
4	0	0	40	Cuba	0

```
[122]: #look at statistics of our numerical data. Validate data makes sense in context.
data.describe()
```

```
[122]:
```

	age	capital-gain	capital-loss	hours-per-week	output
count	32561.000000	32561.000000	32561.000000	32561.000000	32561.000000
mean	38.581647	1077.648844	87.303830	40.437456	0.240810
std	13.640433	7385.292085	402.960219	12.347429	0.427581
min	17.000000	0.000000	0.000000	1.000000	0.000000
25%	28.000000	0.000000	0.000000	40.000000	0.000000
50%	37.000000	0.000000	0.000000	40.000000	0.000000
75%	48.000000	0.000000	0.000000	45.000000	0.000000
max	90.000000	99999.000000	4356.000000	99.000000	1.000000

```
[15]: #rename class to default to avoid python class initiation
data.rename(columns = {'class' : 'output'}, inplace= True)

#Drop education-num as it is the same as education/drop fnlwgt as it is not
↳relevant
data.drop(columns=['education-num', 'fnlwgt'],inplace = True)
```

```
[16]: #Change Target values for visualizations to >50k and <50k
data['output'] = data['output'].replace({0: '<50k'})
data['output'] = data['output'].replace({1: '>50k'})

#Target distribution
data['output'].value_counts()
```

```
[16]: <50k    24720
      >50k     7841
      Name: output, dtype: int64
```

```
[17]: #Analyze missing data
total = data.isnull().sum().sort_values(ascending=False)
percent = (data.isnull().sum()/data.isnull().count()).
↳sort_values(ascending=False)
missing_data = pd.concat([total, percent], axis=1, keys=['Total', 'Percent'])
missing_data.head(20)
```

```
[17]:
```

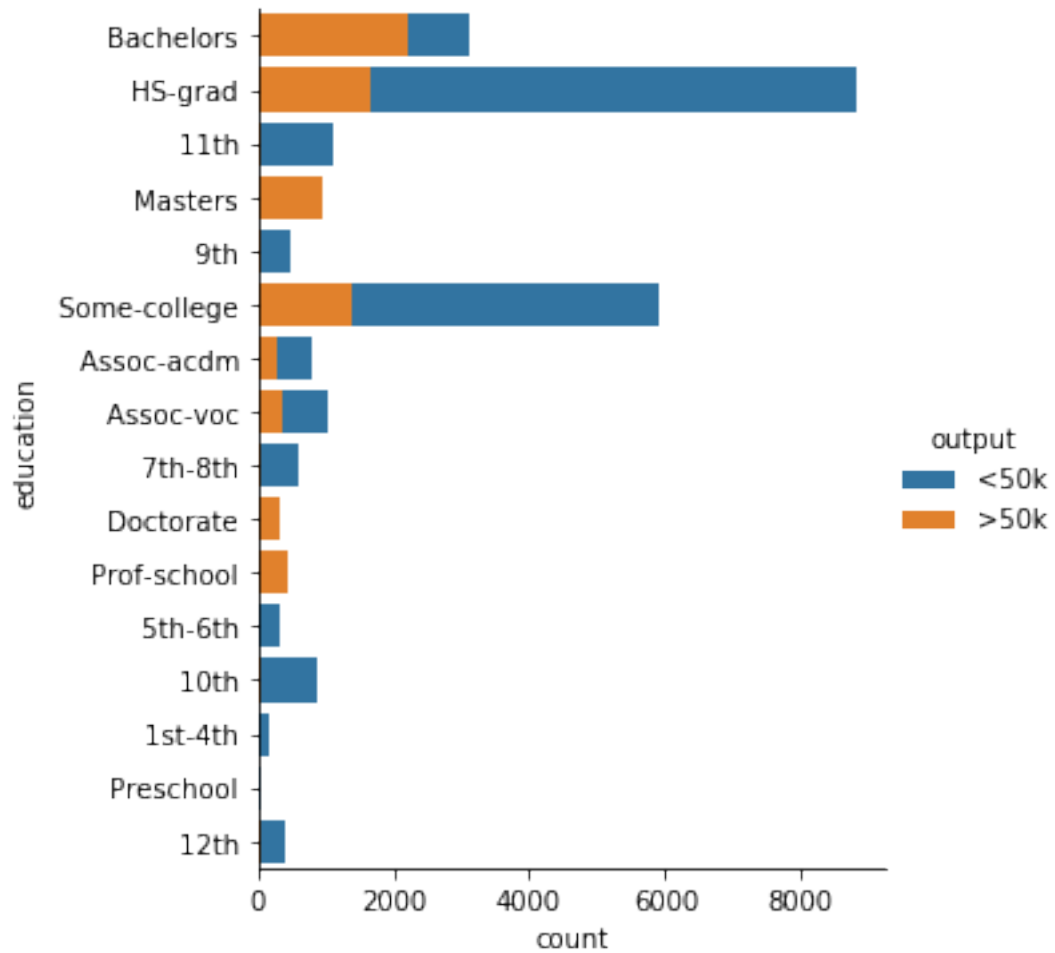
	Total	Percent
native-country	583	0.017905

output	0	0.000000
hours-per-week	0	0.000000
capital-loss	0	0.000000
capital-gain	0	0.000000
sex	0	0.000000
race	0	0.000000
relationship	0	0.000000
occupation	0	0.000000
marital-status	0	0.000000
education	0	0.000000
workclass	0	0.000000
age	0	0.000000


```
[18]: #Delete missing data where observations of na > 1.  
#In this case we are dropping rows where native-country is na  
data = data.drop((missing_data[missing_data['Total'] > 1]).index,1)  
  
#Check to make sure there is no missing data  
data.isnull().sum().max()
```

[18]: 0

```
[19]: #Visualize Target variable by Education Type  
var = 'education'  
ax = sns.catplot(y = var, data=data, kind="count",hue='output',dodge=False)
```




```
[20]: #Visualize Target variable by Sex
var = 'sex'
ax = sns.catplot(y = var, data=data, kind="count", hue='output', dodge=False)
```



```
[95]: #Correlation matrix table

k = 10 #number of variables for heatmap
corrmat = data.corr()
cols = corrmat.nlargest(k, 'output')['output'].index
cm = np.corrcoef(data[cols].values.T)
sns.set(font_scale=1.25)
hm = sns.heatmap(cm, cbar=True, annot=True, square=True, fmt='.2f',
    ↳annot_kws={'size': 10}, yticklabels=cols.values, xticklabels=cols.values)
plt.show()
```



```
[21]: #Change target back to 0 and 1 before modeling
```

```
data['output'] = data['output'].replace({'<50k':0})
data['output'] = data['output'].replace({'>50k':1})
```

```
[22]: #Setup Pycaret
```

```
ctl = setup(data=data, target = 'output',use_gpu = True, normalize=True )
```

Setup Succesfully Completed!

<pandas.io.formats.style.Styler at 0x1e6c7010d48>

```
[23]: #Run 14 Classification models on 10 fold cross validation
```

```
compare_models(fold=10)
```

```
<pandas.io.formats.style.Styler at 0x1e6c5c6b208>
```


```
[23]: <catboost.core.CatBoostClassifier at 0x1e6c410d208>
```

```
[24]: #We decide to go with the Catboost Classifier as it maximizes AUC and F1.  
      CBR = create_model('catboost')
```

```
<pandas.io.formats.style.Styler at 0x1e6c5da4a88>
```

```
[32]: ##SHAP graph helps up understand feature importance of the CATmodel  
      interpret_model(CBR)
```

```
INFO:logs:Initializing interpret_model()  
INFO:logs:interpret_model(estimator=<catboost.core.CatBoostClassifier object at  
0x000001E6C7006148>, plot=summary, feature=None, observation=None)  
INFO:logs:Checking exceptions  
INFO:logs:Importing libraries  
INFO:logs:plot type: summary  
INFO:logs:model type detected: type 2  
INFO:logs:Creating TreeExplainer  
INFO:logs:Compiling shap values
```



INFO:logs:Visual Rendered Successfully
 INFO:logs:interpret_model() succesfully
 completed...

[119]: *#Load in test Data and preprocess for prediction accuracy*

```
test = pd.read_csv('au_test.csv')
test.drop(columns=['education-num', 'fnlwgt'],inplace = True)
test.dropna(how='all',inplace=True)

pred = predict_model(CBR, data=test, verbose=True)
```

[119]: 0.871

[120]: *#The y_pred was in string format. y_pred is converted to string to produce_*
↪accuracy score

```
pred['class'].value_counts()
pred['Label'].value_counts()
```



```
y_test = pred['class'].to_numpy()
y_pred = pred['Label'].to_numpy()
y_pred = [int(i) for i in y_pred]

round(accuracy_score(y_test,y_pred),3)
```

[120]: 0.871

```
[121]: #Confusion Matrix
confusion_matrix(y_test,y_pred)
```

[121]: array([[11684, 751],
 [1347, 2499]], dtype=int64)

```
[ ]: #Save model as a pkl file for later use.
save_model(CBR, 'CBR')
```

Overall, the PyCaret packages combined with some of the sklearn packages works very effectively. We were able to complete an end to end machine learning project in a fraction of the time and code. Pycaret also has packages to deploy model a production setting. For this project we didn't tune the hyperparemeters, but PyCaret has the ability to tune in 1 line of code.